Ethan Ng

Ms. Valenza

Creative Writing, period 11

6 March 2023

Introduction

This research paper delves into the concept of a technological singularity and explores its various facets. This paper aims to provide an overview of what a technological singularity is, consciousness in computer programs and its role in catalyzing a singularity, multiple scenarios on how it could occur, and arguments for it not occurring. This paper will also discuss technological acceleration extrapolated from past singularities and revolutions and how this data will predict if and when a singularity will occur. Through a comprehensive analysis of these factors, the paper aims to provide insight into the concept of singularity and its possibility of happening.

Defining a Singularity

A technological singularity is defined as the point at which the capabilities of artificial intelligence surpass those of human intelligence, leading to extreme technological acceleration beyond our current understanding. As Luchianchicova suggests, “it represents the superiority of machine intelligence over the human person who created it” (Luchianchicova). The singularity represents an absolute and impenetrable limit to progress, or what Kruger calls the last frontier. The metaphor of the final intellectual frontier, as conceptualized by Vinge, Yudkowsky, and Kurzweil, also suggests that the singularity is an impenetrable and insurmountable event horizon for humans, marking “the beginning of an unlimited expansion into the universe, in which

humans are also allowed to participate" (Kruger). The following stages are identified by Kurzweil as the five stages leading up to the realization of the singularity: "1. the origin of matter; 2. the origin of life; 3. the origin of brains/mind; 4. the origin of technology; and 5. the fusion of human and machine intelligence" (Kruger). In Kurzweil's mind, the most probable type of singularity is one where humans and machines meld, allowing humans to transcend their physical and mental limitations.

Defining Consciousness Within Computer Programs

Defining consciousness within computer programs is a complex topic that has been widely debated by researchers and philosophers. Multiple tests have been created to test for intelligence and consciousness in computer programs. Proposed by Alan Turing, the Turing test is a well-known test for determining whether a computer program can exhibit intelligent behavior that is indistinguishable from human behavior. If a human, while conversing with a machine, is convinced that they are actually conversing with a human, the AI has passed the Turing test. A thought experiment, The Chinese Room Argument, proposed by philosopher John Searle in 1980, aims to address the notion of consciousness within computer programs. "Imagine a native English speaker who knows no Chinese locked in a room full of boxes of Chinese symbols (a data base) together with a book of instructions for manipulating the symbols (the program). Imagine that people outside the room send in other Chinese symbols which, unknown to the person in the room, are questions in Chinese (the input). And imagine that by following the instructions in the program the man in the room is able to pass out Chinese symbols which are correct answers to the questions (the output). The program enables the person in the room to pass the Turing test for understanding Chinese but he does not understand a word of Chinese" (Cole). This thought experiment aims to prove that even if a computer program could pass the

Turing test, it would not necessarily have consciousness. It states that passing this test only shows the ability to manipulate symbols, not a true understanding of meaning or consciousness. “The narrow conclusion of the argument is that programming a digital computer may make it appear to understand language but could not produce real understanding. Hence the ‘Turing Test’ is inadequate” (Cole). The Chinese Room Argument implies that even though AI can perform great feats, it does not contain true intelligence and simply echoes data produced solely by humans. However, there are also researchers who argue that consciousness can emerge from complex information processing in computer programs, citing examples such as neural networks and artificial general intelligence. Despite ongoing debates, the question of whether a computer program can possess consciousness remains unresolved and exists as a vital component to determining whether a singularity can or cannot occur.

Technological Acceleration and Its Role in Predicting a Singularity

Our past technological advancements are a measure that many researchers have utilized to predict the next revolution. They factor in events like the renewable energy revolution, the biotechnology revolution, the information age, the industrial revolution, and the agricultural revolution. From plotting this data, researchers have concluded that at “the first stage of the Cybernetic revolution, the speed of technical progress accelerated, and in the second stage (which we have been in since the 1990s), it slowed. We believe that this deceleration will not change until the mid-2030s or beginning of the 2040s, after which technological growth will experience a new acceleration. There will then be a gradual slow-down up to the point of

singularity, with a subsequent change of the pattern" (Grinin et al.).

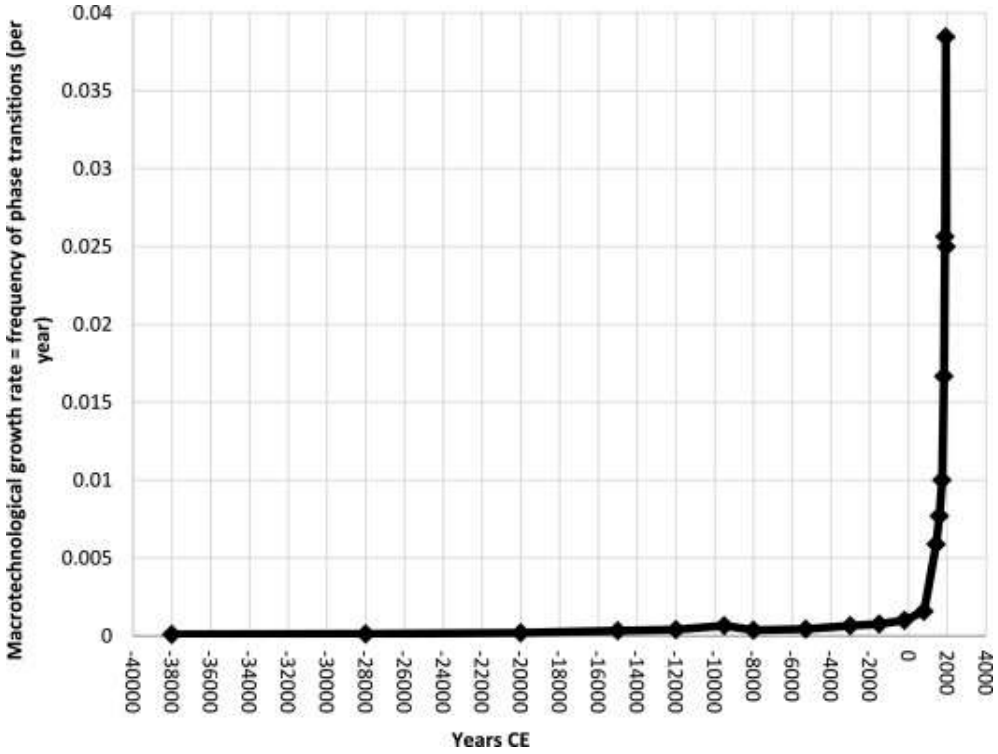

Figure. 1. Macro Technological Growth Rate Graph., Grinin, Leonid, et al. "A Quantitative Analysis of Worldwide Long-term Technology Growth: From 40,000 BCE to the Early 22nd Century." Technological Forecasting and Social Change, vol. 155, June 2020, p. 119955. ScienceDirect.

From extrapolating the technological acceleration data (see figure 1), Grinin et al. theorize that there will be a new acceleration in technological progress. This acceleration will not change until the middle or beginning of the 2040s, at which point they anticipate a fresh acceleration in technological progress. After that, there will be a progressive slowdown until the singularity, after which the pattern will alter. Though we have identified the predicted points of highest acceleration, this does not necessarily mean a singularity will occur. Magee and Devezas argue that "rapid and accelerating technical change can result in profound societal change, but it is not

at all necessary or expected that these societal changes will occur at the time of most rapid improvement in capability" (Magee and Devezas). It is also important to note that an AI singularity, as was defined above, differs tremendously from our past revolutions. We could see an AI revolution where society heavily integrates AI into our lives rather than the predictions of AI dominance or biological fusion.

Ways a Singularity Can Occur

Many ways a singularity could occur have been proposed by various different researchers. Goertzel states that one possibility is that narrow AI (AI designed for a specific task like facial recognition or writing code) gradually decreases in narrowness, a trend we have been seeing recently. It would then turn into AGI – artificial general intelligence. Goertzel claims, "Explicit AGI research doesn't really get anywhere till narrow AI has built up a lot more knowledge about how to solve particular sorts of problems using specialized algorithms" (Goertzel). In the future, narrow AI may become sufficiently intelligent to reach human-level AGI, at which point "AI tech at various levels of generality is already thoroughly integrated into human society, and potentially even into human organisms via brain-computer interfacing technology" (Goertzel). This abundant integration of AI would make sparking a singularity much more likely. Another possibility is that we are currently unknowingly living through a singularity. Famous AI researcher Ray Kurzweil predicts that "AGI is achieved via scanning human brains, figuring out the nature of human thought from these scans, and then replicating human brain function on massively powerful computer hardware (whose existence is predicted by Moore's Law and its brethren). Furthermore, he views this human-level AGI as integrating richly with human society, so that it's not an 'Us versus Them' type scenario, but rather a

scenario in which the boundary between Us and Them is fluid, evolving and impossible to define" (Goertzel). The way AI functions is quite comparable to that of the brain – consisting of layers of nodes that are analogous to neurons in the brain. These two disciplines complement each other, and collaboration in these fields could possibly lead to the creation of human-level AGI. Magee and Devezas conclude that "subjective time speeds up until the rate of progress is merely fast for the inhabitants of the social realm experiencing the change, not as superfast as an observer on the outside and particularly from another time might perceive it" (Magee and Devezas). This evidence supports the claim that singularities are not apparent to the people living through them.

Arguments Against a Singularity Occurring

Such as with arguments for a singularity occurring, there are many arguments for it not occurring. Walsh claims that "Silicon has a significant speed advantage over our brain's wetware, and this advantage doubles every two years or so according to Moore's law. But speed alone does not bring increased intelligence" (Walsh). Our current AI improves by being fed more data and prompts, which improves functionality but does not result in genuine human intelligence. Another claim is that we may be able to develop artificial intelligence that surpasses human intelligence but cannot spark a singularity. Walsh states, "Even if we have enough intelligence to design super-human artificial intelligence, the result may not be adequate to precipitate a technological singularity. Improving intelligence is far harder than just being intelligent" (Walsh). Walsh also claims that "There are many computational problems for which even exponential improvements are not enough to help us solve them practically" (Walsh). For example, the Halting Problem proposed by Alan Turing states that it is impossible to know

whether a given segment of code will stop running or run forever. Some problems, despite how intelligent a system attempting to solve them may be, are impossible to solve, and somewhere along the development of AGI, we may encounter such problems.

Conclusion

The technological singularity remains a topic of scientific debate that has intrigued researchers and scholars for decades. Although the possibility and necessity of sentient AI to catalyze a singularity are unclear, they serve as a useful measure of our technological trajectory. The most reliable method of forecasting a singularity is by quantifying technological acceleration, but this phenomenon is not a reliable predictor of a singularity. Moreover, the acceleration may obfuscate the perception of individuals living through a singularity, making it challenging to identify whether one is happening. Researchers have attempted to leverage past technological advancements to predict the next revolution, but the probability of a singularity occurring remains uncertain. Most theories are speculative in nature, with little empirical data available to either support or reject the likelihood of a singularity. More research is required to draw definitive conclusions about the potentiality and timeline of a singularity.

Works Cited

Cole, David. "The Chinese Room Argument." *The Stanford Encyclopedia of Philosophy*, winter 2020, plato.stanford.edu/archives/win2020/entries/chinese-room/. Accessed 7 Feb. 2023.

Goertzel, Ben. "Human-level Artificial General Intelligence and the Possibility of a Technological Singularity: A Reaction to Ray Kurzweil's the Singularity Is Near, and McDermott's Critique of Kurzweil." *Artificial Intelligence*, vol. 171, no. 18, Dec. 2007, pp. 1161-73. *Academic Search Index*, https://doi.org/10.1016/j.artint.2007.10.011.

Grinin, Leonid, et al. "A Quantitative Analysis of Worldwide Long-term Technology Growth: From 40,000 BCE to the Early 22nd Century." *Technological Forecasting and Social Change*, vol. 155, June 2020, p. 119955. *ScienceDirect*, https://doi.org/10.1016/j.techfore.2020.119955.

Kruger, Oliver. "'The Singularity Is Near!' Visions of Artificial Intelligence in Posthumanism and Transhumanism." *International Journal of Interactive Multimedia and Artificial Intelligence*, vol. 7, nos. 1 SI, Sept. 2021, p. 16+. *Gale Academic OneFile*, link.gale.com/apps/doc/A676106453/AONE?u=nysl_li_jhsch&sid=ebsco&xid=b7393cc8. Accessed 10 Feb. 2023.

Luchianchicova, Miroslava Metleaeva. "Artificial Intelligence, Transhumanism and Translation as a Universal Property of Thinking." *Bulletin of Integrative Psychiatry*, vol. 27, no. 3, Sept. 2021, p. 109+. *Gale Academic OneFile*, link.gale.com/apps/doc/A678979875/AONE?u=bocesnas12&sid=googleScholar&xid=10ba44d5. Accessed 14 Feb. 2023.

Magee, Christopher L., and Tessaleno C. Devezas. "How Many Singularities Are near and How Will They Disrupt Human History?" *Technological Forecasting and Social Change*, vol.

78, no. 8, Oct. 2011, pp. 1365-78. *ScienceDirect*, https://doi.org/10.1016/j.techfore.2011.07.013. Accessed 8 Feb. 2023.

Walsh, Toby. "What If... We Create Human-Level Artificial Intelligence?" *New Scientist*, vol. 232, no. 3100, 19 Nov. 2016, pp. 32-34. *Academic Search Elite*, https://doi.org/10.1016/S0262-4079(16)32126-1.